

# Solving Nonstationary Non-Markovian Queueing Systems

A Transfer Learning Neural Network Approach

Yunan Liu

Principal Research Scientist, Supply Chain Optimization Technology

Joint work with Spyros Garyfallos, Amazon Connect, AWS

AMLC 2025

# Queueing Theory



Long waiting line at the AMLC registration desk.

- A mathematical study of waiting lines or queues (a branch of operations research).
- Applications: customer call centers, healthcare, manufacturing systems, transportation, cloud computing, food services, etc.
- Entities in a queueing system:
  - ▶ Customers: consumers, patients, jobs, data packets, riders, etc.
  - ▶ Servers: agents, doctors, nurses, beds, machines, cloud, drivers, etc.

# Real-World Service Systems

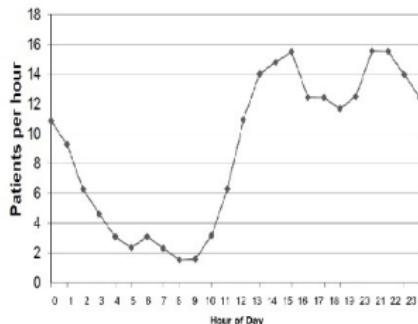
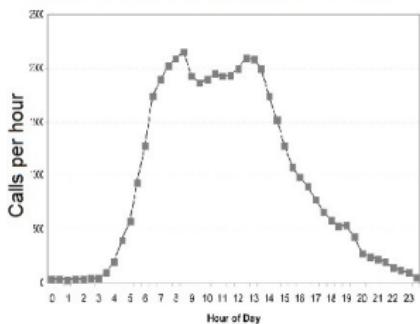
## Complex dynamics:

- **Time-varying demand:** Customer arrivals fluctuate throughout the day
- **Time-varying capacity:** Staffing levels adjust to match demand
- **Non-Markovian structure:** Service and abandonment times don't follow exponential distributions
- **Multiple performance objectives:** Need to satisfy concurrent service-level constraints
- **Others:** Complex network topologies, multiple customer classes, customer/server behavior.

## Key Question:

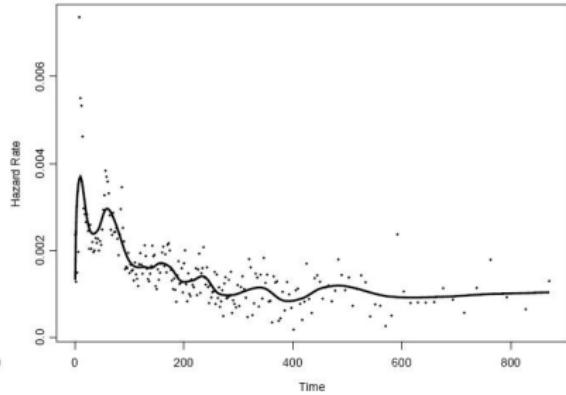
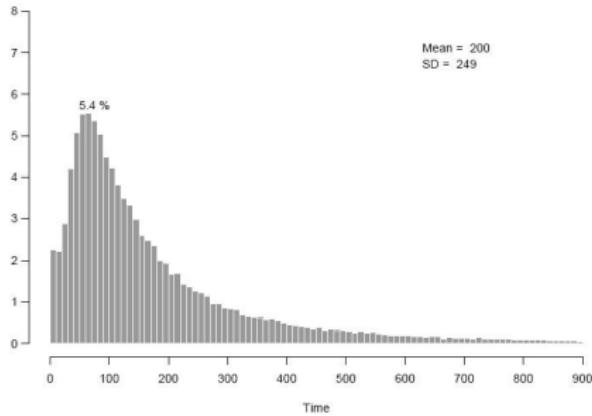
- **Performance prediction:** *How can we accurately predict system performance in such complex systems?*
- **Capacity planning:** *How can we optimally plan service capacity to meet service-level targets under nonstationary non-Markovian dynamics?*
- **Dynamic control:** *What real-time actions should be taken when service-level constraints are at risk of being violated?*

# Time-Varying Demand



- **Call centers:** Peak hours during business days, lunch breaks
- **Healthcare:** Emergency department arrivals follow daily patterns
- **Cloud services:** Traffic varies by time of day, day of week

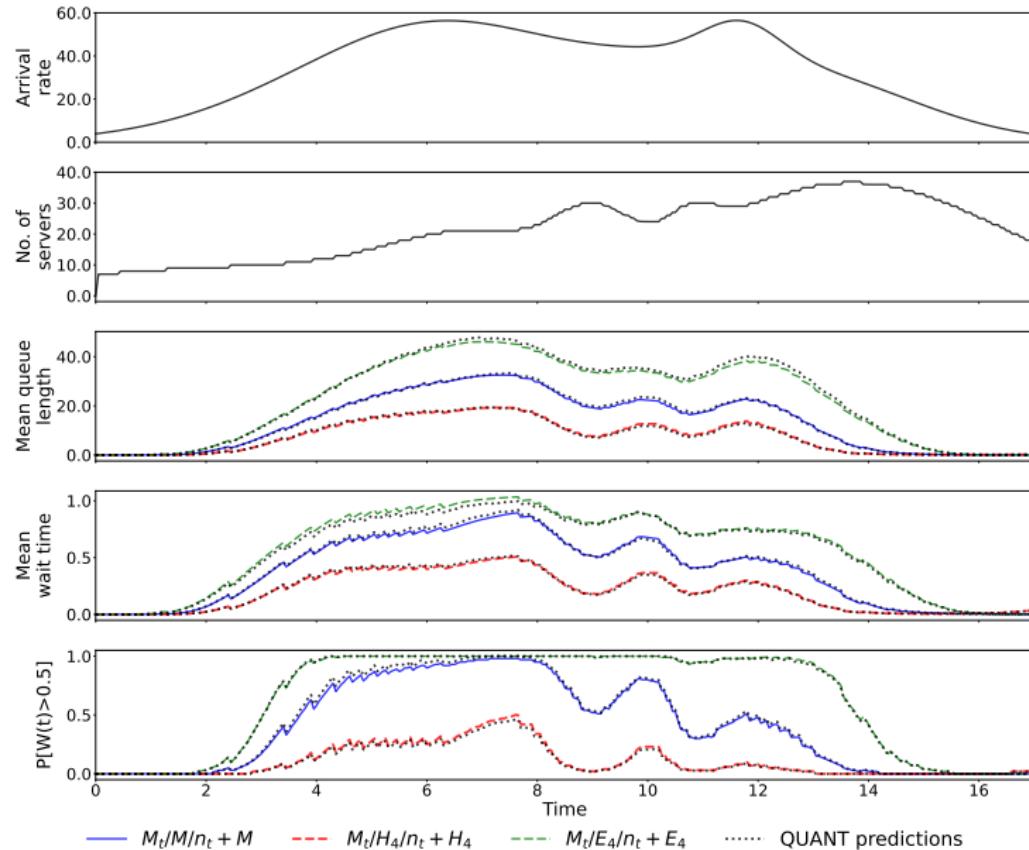
# Non-Markovian Distributions



**Real service and abandonment times are NOT exponential:**

- Call durations often have high variability (lognormal)
- Customer abandonment times show complex patterns
- Distribution shape significantly impacts performance

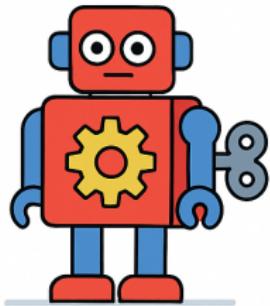
# Markovian vs. Non-Markovian: Performance Divergence



# Markovian Model vs. Non-Markovian Model

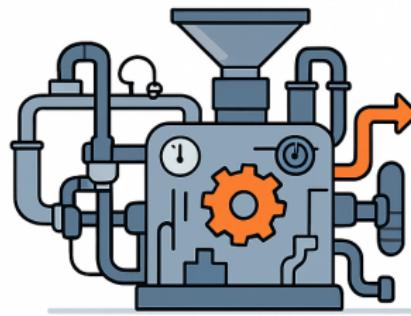
**Methodological complexity:**

**Markovian world**



**A Simple Toy**

**Non-Markovian world**



**A Complex Machine**

# Markovian Model vs. Non-Markovian Model

## Markovian queueing systems:

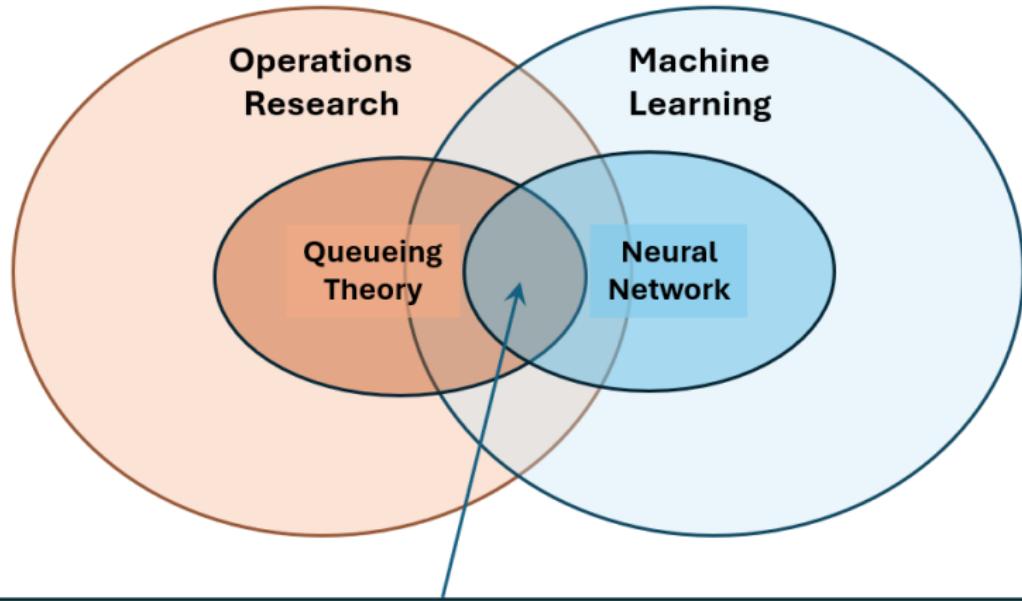
- Memoryless property: Future depends only on current state
- State = simple queue length
- Exact solutions via uniformization and matrix methods
- Computationally tractable

## Non-Markovian queueing systems:

- Must track customer ages, residual service times
- State space becomes infinite-dimensional
- No closed-form solutions
- Heavy-traffic approximations are complex and often inaccurate
- Simulation is slow and computationally expensive

# QUANT: Q**U**eueing Analysis via Neural-Network Transfer Learning

# Where Are We?



**QUANT – QUeueing Analysis via Neural-network Transfer learning**

# The $G_t/GI/n_t + GI$ Queueing Model

## System Components:

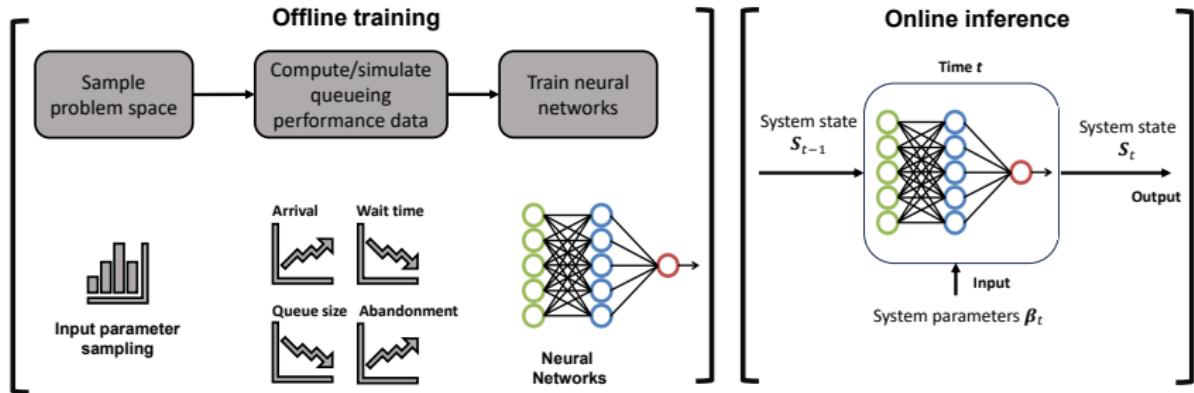
- $G_t$ : General arrival process with nonstationary rate  $\lambda(t)$
- $GI$ : General independent service times (distribution  $G$ , mean  $1/\mu$ )
- $n_t$ : Time-varying number of servers  $n(t)$
- $+GI$ : Customer abandonment (distribution  $F$ , mean  $1/\theta$ )

## Performance Metrics to Predict:

- Mean waiting time:  $\mathbb{E}[W(t)]$
- Mean queue length:  $\mathbb{E}[Q(t)]$
- Mean busy servers:  $\mathbb{E}[B(t)]$
- Tail probability of delay:  $\mathbb{P}(W(t) \leq \tau)$  (e.g.,  $\tau = 60$  sec)

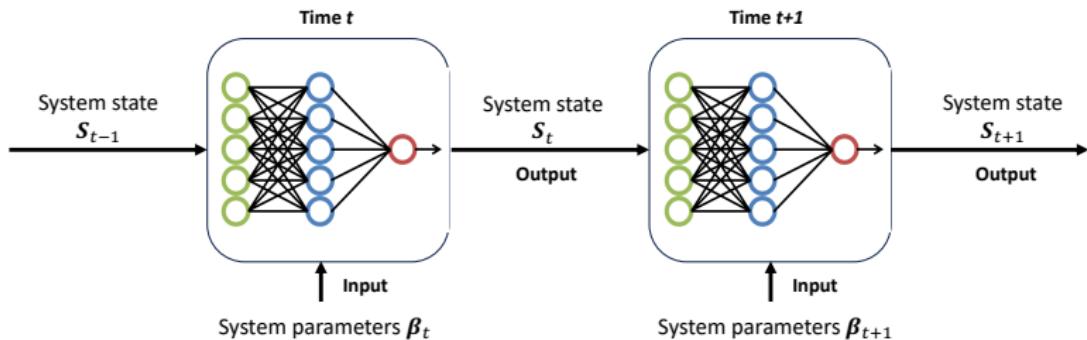
**Goal:** Predict these metrics over entire horizon  $t \in [0, T]$  for any  $\lambda(t)$  and  $n(t)$

# The QUANT Framework



- Goal: Predicts queue performance under time-varying demand and capacity
- Method: Neural network models trained on exact and simulated data
  - ▶ Offline phase: generate data and train networks
  - ▶ Online phase: fast prediction in real time
  - ▶ Inputs: demand rate, staffing level, current system state
  - ▶ Outputs: waiting time, queue length, busy servers, service level.

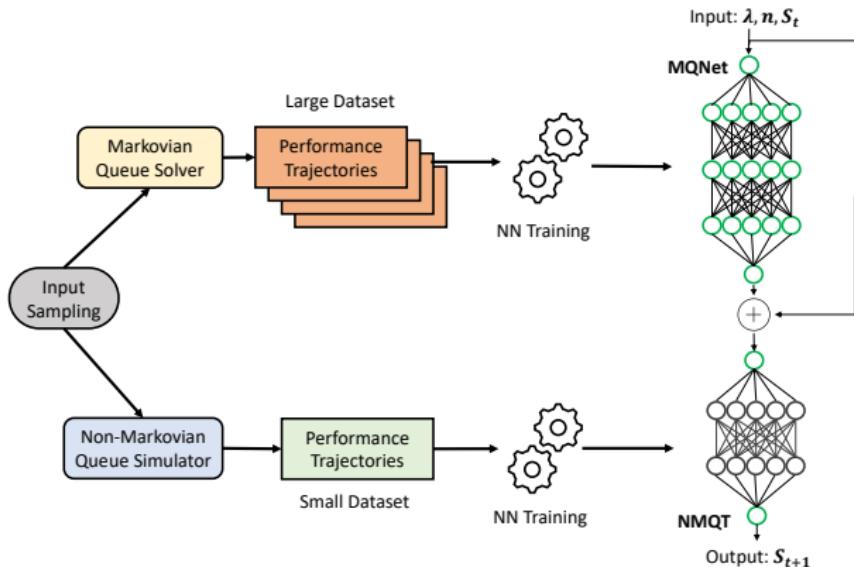
# QUANT Architecture: Macroscopic Level



## Recurrent Structure:

- One-step prediction:  $\mathbf{S}_{t+1} = \mathcal{F}(\mathbf{S}_t, \beta_{t+1})$
- State:  $\mathbf{S}_t = (\mathbb{E}[W(t)], \mathbb{E}[Q(t)], \mathbb{E}[B(t)])$
- Input:  $\beta_{t+1} = (\lambda(t+1), n(t+1))$
- Roll out predictions recursively over time horizon
- Each step takes  $\sim 1$  millisecond on GPU

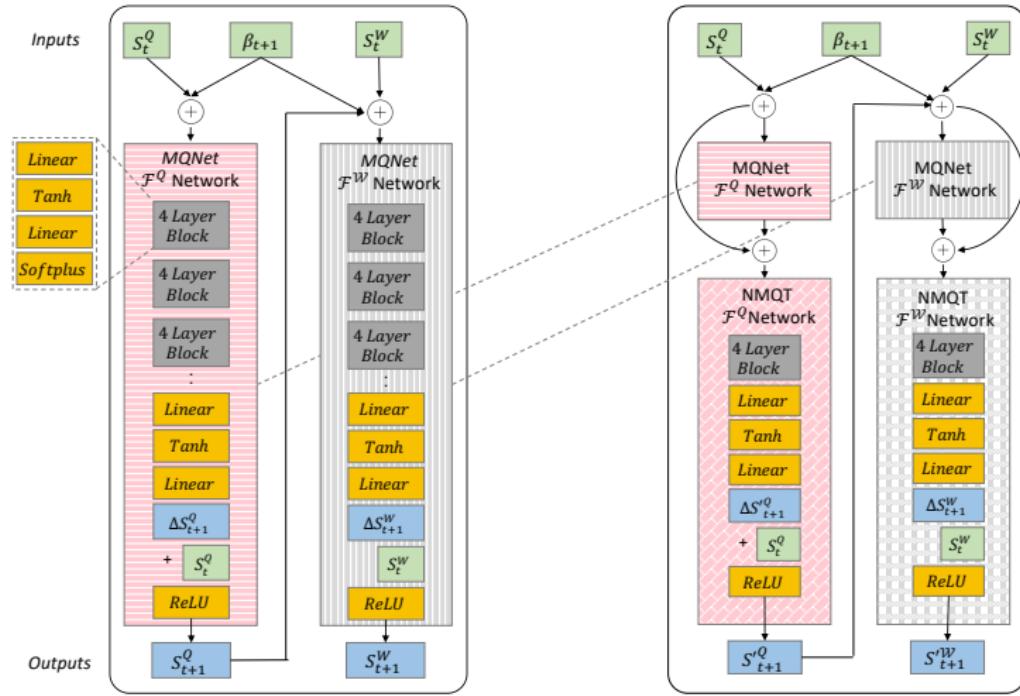
# QUANT Architecture: Mesoscopic Level



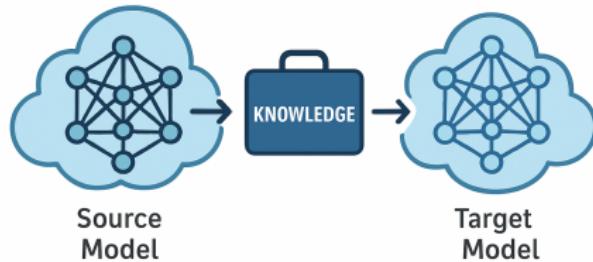
## Two-Module Design:

- **Markovian Queue Network (MQNet):** Pre-trained on  $M_t/M/n_t + M$  systems (exact results)
- **Non-Markovian Queue Transformer (NMQT):** Learns corrections for non-Markovian  $G_t/G/n_t + GI$  (simulated results)

# QUANT Architecture: Microscopic Level



# Why Transfer Learning?



# Features Characterizing A Queueing System

## 1. System Structure

- Topology: single, pooled, parallel, tandem, or networked queues
- Capacity: number of servers; time-varying or fixed

## 2. Operational Rules

- Service discipline: FCFS, LCFS, Priority, Processor Sharing
- Routing: fastest-server-first, join-shortest-queue
- Customer behavior: balking, abandonment, retrial

## 3. Random Elements: Averages

- Arrival rate: constant or time-varying
- Average service time, average abandonment time

## 4. Random Elements: Distribution beyond Means

- Arrival process: Poisson vs. non-Poisson
- Service- and abandonment-time distribution: exponential vs. general

# Features Characterizing A Queueing System

## 1. System Structure

- Topology: single, pooled, parallel, tandem, or networked queues
- Capacity: number of servers; time-varying or fixed

## 2. Operational Rules

- Service discipline: FCFS, LCFS, Priority, Processor Sharing
- Routing: fastest-server-first, join-shortest-queue
- Customer behavior: balking, abandonment, retrial

## 3. Random Elements: Averages

- Arrival rate: constant or time-varying
- Average service time, average abandonment time

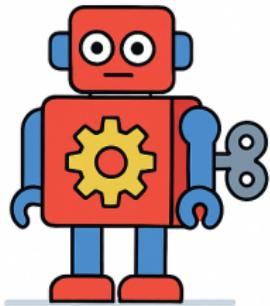
## 4. Random Elements: Distribution beyond Means

- Arrival process: Poisson vs. non-Poisson
- Service- and abandonment-time distribution: exponential vs. general

# Markovian Model vs. Non-Markovian Model

**Methodological complexity:**

**Markovian world**



**A Simple Toy**

**Non-Markovian world**



**A Complex Machine**

# Transfer Learning: The Key Innovation

## Why Transfer Learning?

- **Markovian systems:** Exact solutions available via MQSolver (cheap!)
- **Non-Markovian systems:** Only simulation available (expensive!)
- Generate 10,000 trajectories for Markovian training:  $\sim 7$  CPU-days
- Generate 10,000 trajectories for non-Markovian:  $\sim 100$  CPU-days

## Our Solution:

- ① Pre-train MQNet on 10,000 Markovian trajectories
- ② Fine-tune with only 100 non-Markovian trajectories (1% of data!)
- ③ Achieves near-optimal accuracy with 99% less simulation
- ④ Total training:  $\sim 8$  CPU-days vs.  $\sim 100$  CPU-days

⇒ **Orders of magnitude reduction in training cost**

# Data Generation Strategy

## Synthetic Trajectory Generation:

- Time-series composition of Gaussian processes
- Creates realistic nonstationary patterns: single-peak, multi-peak, high-volatility
- Covers operational bounds:  $\lambda \in [0, 100]$ ,  $n \in [1, 100]$

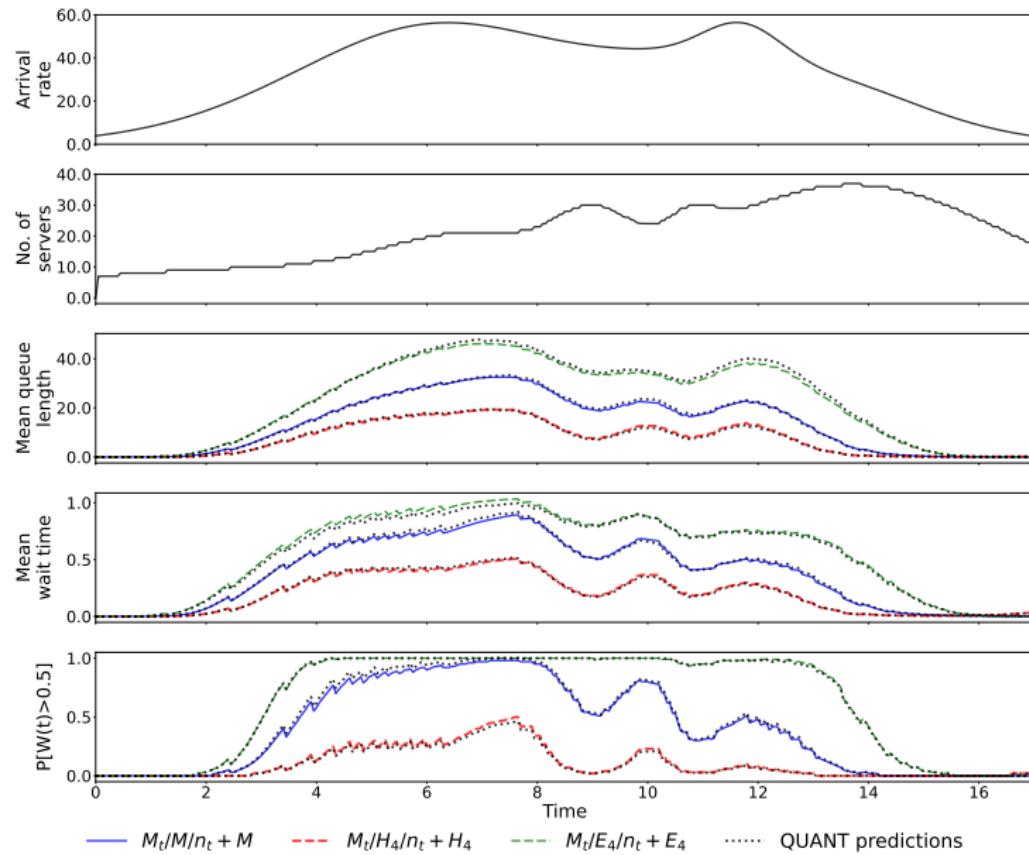
## MQSolver: Exact Markovian Solution

- Exploits birth-death process structure
- Uniformization for time-inhomogeneous systems
- Computes exact  $\mathbb{P}(W(t) \leq \tau)$  for all  $t, \tau$
- Runtime: 60 seconds per trajectory (vs. 900s for simulation)

## Non-Markovian Simulation:

- 10,000 replications for smooth estimates
- Only 100 trajectories needed for transfer learning

# QUANT Prediction: An Example



# Runtime Performance

Method	Runtime (seconds)
QUANT (GPU)	<b>0.200</b>
QUANT (CPU)	0.350
MQSolver	60
Simulation	900

## Speedup Analysis:

- **300× faster** than exact Markovian solver
- **4,500× faster** than simulation
- Enables real-time optimization with thousands of evaluations
- Critical for production capacity planning systems

# Prediction Accuracy and Robustness

Configuration	TAMSE <sup>1</sup>
<b>Base Markovian System</b>	
Queue/Busy metrics	0.09999
Waiting time	0.00131
End-to-end	<b>0.07032</b>
<b>By System Scale</b>	
Small scale (0-41 servers)	0.08621
Medium scale (41-72 servers)	0.07559
Large scale (72-100 servers)	0.06377
<b>By Distribution (Non-Markovian)</b>	
$M_t/H_2/n_t + H_2$ (SCV=4)	0.07144
$M_t/E_4/n_t + E_4$ (SCV=1/4)	0.07092

## Key Findings:

- Consistent accuracy (error < 0.1) across all configurations
- Minimal degradation from Markovian to non-Markovian
- Scale-invariant performance

---

<sup>1</sup>TAMSE: time-averaged mean square error

# Transfer Learning Efficiency

## How much non-Markovian data is actually needed?

Fine-tuning Data	TAMSE <sup>2</sup>	Relative Error
Baseline (10,000 trajectories, no transfer)	0.0703	–
50% (5,000)	0.0709	0.8%
10% (1,000)	0.0714	1.6%
<b>1% (100)</b>	<b>0.0734</b>	<b>4.4%</b>
0.5% (50)	0.0872	24.0%

### Key Insight:

- Robust performance down to **1% of data (100 trajectories)**
- Accuracy within 5% of full dataset
- Sharp degradation below this threshold
- **Practical guideline: Use 100-500 trajectories for new systems**

---

<sup>2</sup>TAMSE: time-averaged mean square error

# QUANT-based Optimal Staffing

# Optimal Staffing under Service-Level Constraints

**Problem:** Given  $\lambda(t)$  and constraints, find optimal  $n^*(t)$

**Example Constraints:**

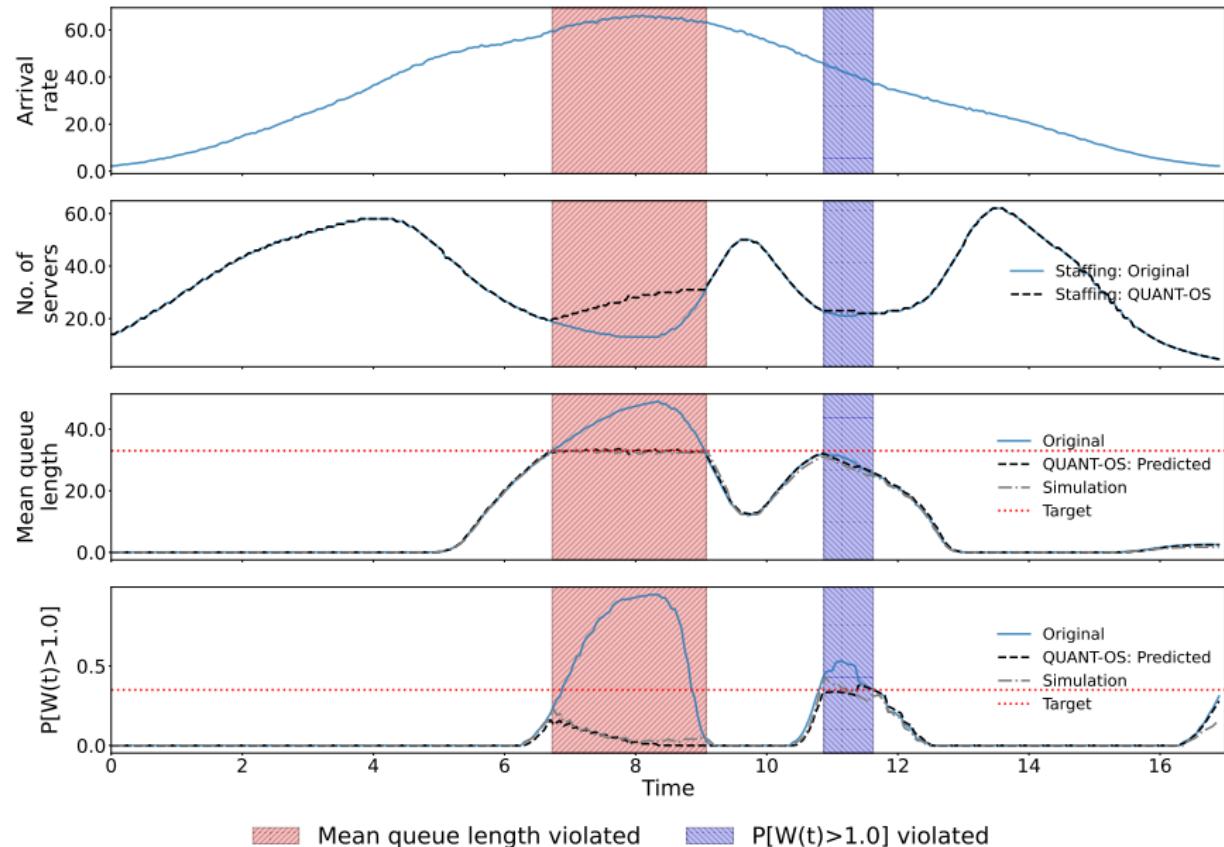
- Mean waiting time:  $\mathbb{E}[W(t)] \leq w$  (e.g.,  $w = 30$  seconds)
- Mean queue length:  $\mathbb{E}[Q(t)] \leq q$  (e.g.,  $q = 10$ )
- Tail probability:  $\mathbb{P}(W(t) > \tau) \leq \alpha$  (e.g.,  $\tau = 60$  seconds,  $\alpha = 20\%$ )
- Meeting multiple constraints simultaneously!

**Three Optimization Methods:**

- ① **Binary Search:** Exploits monotonicity,  $\sim 8$  iterations
- ② **Batch Evaluation:** Parallel GPU evaluation, 1 pass
- ③ **Gradient-Based:** Auto-differentiation through QUANT,  $\sim 35$  iterations

All methods complete in seconds (vs. hours for simulation-based optimization)

# QUANT-OS: A Multi-Constraint Example



# Optimization Performance

Method	Iterations	Time (s)	Success
Binary Search	8.2	1.2	100%
Batch Evaluation	1.0	0.8	100%
Gradient-Based	34.5	10.1	97.2%

## Validation Against Simulation:

Constraint Type	Pred. Error	Violation Rate
$\mathbb{E}[W] \leq w$	0.082-0.107	2.5-3.5%
$\mathbb{E}[Q] \leq q$	0.080-0.096	2.0-3.0%
$\mathbb{P}(W > \tau) \leq \alpha$	0.131	5.0%
Multi-constraint	0.140	6.0%

⇒ Over 94% constraint satisfaction in all cases

# Robustness Across Configurations

Configuration	Convergence	Satisfaction	Pred. Error
$M_t/M/n_t + M$	99.2%	98.5%	0.121
$M_t/E_4/n_t + E_4$	98.5%	97.0%	0.224
$M_t/H_2/n_t + H_2$	97.8%	95.5%	0.360
Small scale	96.7%	96.5%	0.156
Medium scale	97.9%	97.0%	0.140
Large scale	98.4%	98.5%	0.132

## Key Findings:

- Consistent performance across all distributions tested
- Works for both high and low variability systems
- Scale-independent (works for small and large systems)
- Maintains > 95% satisfaction even for high-variability systems

# Conclusions

## Summary:

- Recurrent neural network for nonstationary non-Markovian queues with time-varying capacity
- Transfer learning achieves 99% data reduction
- Real-time multi-constraint optimization in seconds

## Future directions:

- Multiple customer classes with priority scheduling
- Network topologies using Graph Neural Networks
- Reinforcement learning with QUANT as environment model
- Multi-objective optimization beyond constraint satisfaction

# Thank You!

Spyros Garyfallos: [spyrosg@amazon.com](mailto:spyrosg@amazon.com)

Yunan Liu: [yunanliu@amazon.com](mailto:yunanliu@amazon.com)

Paper: Transfer Learning Neural Networks for  
Nonstationary Non-Markovian Queueing Systems.  
*Submitted to INFORMS Journal on Computing.*

Website: [yliu48.github.io](https://yliu48.github.io)